

LocalGovBench v0.1 — Delphi Round 1 participant questionnaire

Framework version: 0.1

Round: 1

Generated: 2026-06-18

Criteria: 25

Complete one block per criterion. See `participant_instructions.md` for scales.

Legal and Regulatory Compliance (`legal_regulatory`)

1. `legal_regulatory_gdpr_readiness`

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: GDPR readiness for training, inference, and logging data flows.

Criterion description: GDPR readiness for training, inference, and logging data flows.

Documentation hint (for field study context only; not rated in Round 1):

Records of processing, privacy notices, DPIA where applicable, and DPO consultation notes.

Risk if missing (indicative):

Unlawful or undocumented processing; difficulty demonstrating accountability to supervisory authorities.

Traceability references:

- **ART** — Accountability: Demonstrate processing accountability. GDPR accountability principle requires demonstrable compliance for training, inference, and logs.
- **EU Trustworthy AI** — Privacy and data governance: Lawful and fair processing. HLEG trustworthy AI list emphasises data governance alongside fundamental rights.
- **AI Act** — Risk management and documentation: Records and DPIA documentation. Deployer documentation duties intersect with governance records for data-intensive LLM systems.
- **Mittelstadt et al.** — Inconclusive risk: Prevent opaque data use. Undocumented flows obscure harms analogous to informational inconclusiveness in socio-technical systems.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

2. `legal_regulatory_ai_act_alignment`

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: EU AI Act alignment documented for the deployer role and use case.

Criterion description: EU AI Act alignment documented for the deployer role and use case.

Documentation hint (for field study context only; not rated in Round 1):

Use-case classification memo, risk management notes, human oversight description, and technical documentation references.

Risk if missing (indicative):

Gaps in deployer obligations and post-deployment monitoring may go unidentified until external audit or incident.

Traceability references:

- **ART** — Responsibility: Align deployer obligations with use case. Assigns institutional responsibility for classifying and governing municipal LLM deployments.
- **EU Trustworthy AI** — Technical robustness and safety: Trustworthy design and deployment. ALTAI prompts assessment of safety and robustness before operational reliance.
- **AI Act** — Risk management: Risk-based governance. AI Act risk management expectations inform internal classification and monitoring design.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

3. legal_regulatory_data_retention

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Data retention and deletion rules defined for prompts, logs, and embeddings.

Criterion description: Data retention and deletion rules defined for prompts, logs, and embeddings.

Documentation hint (for field study context only; not rated in Round 1):

Retention schedule, automated deletion configuration, and backup handling procedures.

Risk if missing (indicative):

Excessive retention of prompts or logs increases exposure and complicates erasure requests.

Traceability references:

- **ART** — Accountability: Limit storage of prompts and logs. Retention rules make processing demonstrable and bounded over time.
- **EU Trustworthy AI** — Privacy and data governance: Data minimisation. Trustworthy AI requires proportionate data use across the lifecycle.
- **AI Act** — Data governance: Storage limitation. Governance of logs and embeddings supports post-market monitoring without excess retention.
- **Mittelstadt et al.** — Unfair outcomes: Reduce harm from stale personal data. Excessive retention increases risk of discriminatory or harmful reuse of interaction data.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

4. legal_regulatory_lawful_basis

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion

relevant and clearly worded for programme-level governance readiness assessment? Criterion: Lawful basis and purpose limitation recorded for each data category used.

Criterion description: Lawful basis and purpose limitation recorded for each data category used.

Documentation hint (for field study context only; not rated in Round 1):

Lawful basis register entries, purpose statements in system design documents, and data minimization checklist.

Risk if missing (indicative):

Processing may lack demonstrable legal grounding; harder to justify secondary use of interaction data.

Traceability references:

- **ART** — Accountability: Document purpose and legal basis. Purpose limitation and lawful basis are core accountability artefacts under GDPR.
- **EU Trustworthy AI** — Privacy and data governance: Lawful processing. HLEG links lawful processing to trustworthy municipal AI.
- **AI Act** — Data governance: Lawfulness of training and inference data. Deployer governance must trace datasets used in local LLM operations.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

5. legal_regulatory_cross_border_avoidance

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Cross-border data transfer avoidance for on-premise workloads.

Criterion description: Cross-border data transfer avoidance for on-premise workloads.

Documentation hint (for field study context only; not rated in Round 1):

Architecture diagram showing EU hosting, egress controls, vendor sub-processor list, and network policy excerpts.

Risk if missing (indicative):

Unintended transfers via APIs, telemetry, or cloud-backed components may undermine sovereignty claims.

Traceability references:

- **ART** — Responsibility: Control data location and egress. Institutional responsibility for sovereignty claims over on-premise workloads.
- **EU Trustworthy AI** — Technical robustness and safety: Resilience and control. Operational control supports resilience when external dependencies are limited.
- **AI Act** — Data governance: Sovereign hosting. Data governance for deployers includes knowing where inference and logs reside.
- **Mittelstadt et al.** — Misguided agency: Avoid hidden external agency. Covert cloud egress can misattribute decisions to local authority control.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	

Field	Your response
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

Technical and Security Readiness (technical_security)

6. technical_security_local_architecture

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Local deployment architecture documented end-to-end.

Criterion description: Local deployment architecture documented end-to-end.

Documentation hint (for field study context only; not rated in Round 1):

Architecture diagrams, component inventory, segmentation model, and secrets management approach.

Risk if missing (indicative):

Opaque dependencies complicate incident response and sovereignty review.

Traceability references:

- **ART** — Transparency: Document end-to-end deployment. Architecture transparency enables oversight of local LLM components.
- **EU Trustworthy AI** — Technical robustness and safety: Technical robustness. HLEG robustness requirement applies to system design documentation.
- **AI Act** — Documentation and record-keeping: Technical documentation. Deployer technical documentation theme for traceable system descriptions.
- **Mittelstadt et al.** — Inconclusive risk: Reduce opaque dependencies. Opaque stacks create inconclusive risk assessment for citizens and auditors.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

7. technical_security_access_control

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Access control enforces least privilege for operators and end users.

Criterion description: Access control enforces least privilege for operators and end users.

Documentation hint (for field study context only; not rated in Round 1):

IAM role matrix, authentication method description, and periodic access review records.

Risk if missing (indicative):

Over-privileged accounts increase risk of data exfiltration or model misuse.

Traceability references:

- **ART** — Responsibility: Enforce least privilege. Assigns responsibility for access decisions affecting model and data misuse.

- **EU Trustworthy AI** — Technical robustness and safety: Security and control. Access control is a baseline trustworthy AI security practice.
- **AI Act** — Risk management: Cybersecurity governance. Security measures form part of deployer risk management for AI systems.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

8. technical_security_logging

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Logging captures security-relevant and operational events without excess personal data.

Criterion description: Logging captures security-relevant and operational events without excess personal data.

Documentation hint (for field study context only; not rated in Round 1):

Logging policy, sample log schemas, and redaction or pseudonymization rules.

Risk if missing (indicative):

Forensics and accountability are weakened after security or safety incidents.

Traceability references:

- **ART** — Accountability: Security and accountability logging. Logs support ex-post accountability without unnecessary personal data.
- **EU Trustworthy AI** — Transparency: Traceability of operations. Operational transparency for auditors and oversight bodies.
- **AI Act** — Post-market monitoring: Monitoring and logging. Logging practices underpin monitoring of system behaviour over time.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

9. technical_security_auditability

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Auditability supports traceability of configuration and inference changes.

Criterion description: Auditability supports traceability of configuration and inference changes.

Documentation hint (for field study context only; not rated in Round 1):

Audit trail configuration, change management tickets, and immutable log store references.

Risk if missing (indicative):

Inability to reconstruct who changed models, prompts, or policies when investigating harm.

Traceability references:

- **ART** — Transparency: Immutable audit trails. Auditability makes configuration and inference changes visible.
- **EU Trustworthy AI** — Transparency: Traceability. Trustworthy AI transparency includes traceable changes.
- **AI Act** — Documentation and record-keeping: Record-keeping. Supports reconstructing deployer actions after incidents.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):**Comment (optional):****10. technical_security_model_updates**

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Model update management defines testing, approval, and rollback.

Criterion description: Model update management defines testing, approval, and rollback.

Documentation hint (for field study context only; not rated in Round 1):

Model change procedure, validation results, version registry, and rollback runbook.

Risk if missing (indicative):

Uncontrolled updates may introduce drift, bias, or safety regressions without notice.

Traceability references:

- **ART** — Responsibility: Controlled model change. Named approval paths for updates assign responsibility for model risk.
- **EU Trustworthy AI** — Technical robustness and safety: Robustness over lifecycle. Lifecycle testing aligns with HLEG safety expectations.
- **AI Act** — Risk management: Risk management of changes. Model updates trigger reassessment under deployer risk management.
- **Mittelstadt et al.** — Inconclusive risk: Prevent unreviewed capability shifts. Sudden model changes without review recreate inconclusive risk for users.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):**Comment (optional):****Organizational Governance (organizational)**

11. organizational_accountability

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Accountability structures link AI outcomes to public service responsibilities.

Criterion description: Accountability structures link AI outcomes to public service responsibilities.

Documentation hint (for field study context only; not rated in Round 1):

Governance charter excerpt, committee terms of reference, and service owner nomination.

Risk if missing (indicative):

Responsibility may be diffused between IT vendors and line managers after failures.

Traceability references:

- **ART** — Accountability: Link AI outcomes to public duties. Core ART construct operationalised for municipal LLM programmes.
- **EU Trustworthy AI** — Human agency and oversight: Organisational governance. Institutional governance enables human agency at organisational level.
- **AI Act** — Risk management: Governance structure. Accountability structures support AI Act governance of deployer duties.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

12. organizational_ownership

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Ownership assigns sustained stewardship for each on-premise LLM system.

Criterion description: Ownership assigns sustained stewardship for each on-premise LLM system.

Documentation hint (for field study context only; not rated in Round 1):

RACI or ownership matrix, service catalogue entries, and escalation contacts.

Risk if missing (indicative):

Systems may become orphaned when staff rotate or projects end.

Traceability references:

- **ART** — Responsibility: Sustained service ownership. Ownership assigns ongoing responsibility beyond project phases.
- **EU Trustworthy AI** — Human agency and oversight: Clear roles. Ownership clarifies who exercises oversight on behalf of the public.
- **AI Act** — Risk management: Governance roles. Risk owners identified for deployer obligations.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

13. organizational_role_definition

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Role definitions cover legal, technical, data, and operational functions.

Criterion description: Role definitions cover legal, technical, data, and operational functions.

Documentation hint (for field study context only; not rated in Round 1):

Role descriptions, training plans, and interdisciplinary workshop minutes.

Risk if missing (indicative):

Critical safeguards may be assumed rather than assigned to named roles.

Traceability references:

- **ART** — Responsibility: Interdisciplinary roles. Defines who is responsible for legal, technical, and operational safeguards.
- **EU Trustworthy AI** — Human agency and oversight: Competence and training. Role clarity supports meaningful staff agency in AI operations.
- **AI Act** — Risk management: Organisational measures. Human resources measures in risk management systems.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

14. organizational_procurement_governance

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Procurement governance addresses AI components, support, and exit.

Criterion description: Procurement governance addresses AI components, support, and exit.

Documentation hint (for field study context only; not rated in Round 1):

Contract clauses on model updates, SLAs, audit rights, and transition assistance.

Risk if missing (indicative):

Vendor lock-in and unclear exit paths can constrain future sovereignty choices.

Traceability references:

- **ART** — Responsibility: Contractual governance of vendors. Procurement assigns responsibility for vendor performance and exit.
- **EU Trustworthy AI** — Diversity, non-discrimination and fairness: Diversity and fairness in supply chain. Fair procurement reduces biased or opaque vendor dependencies.
- **AI Act** — Risk management: Supply chain governance. Governance of third-party components in deployer risk programmes.
- **Mittelstadt et al.** — Unfair outcomes: Avoid lock-in harms. Lock-in can unfairly limit future municipal choices and public value.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

15. organizational_risk_ownership

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Risk ownership assigns identification, treatment, and reporting of AI risks.

Criterion description: Risk ownership assigns identification, treatment, and reporting of AI risks.

Documentation hint (for field study context only; not rated in Round 1):

Risk register entries, risk appetite statement, and periodic review cadence.

Risk if missing (indicative):

Risks may be tracked only as technical issues without public-value framing.

Traceability references:

- **ART** — Accountability: Institutional risk register. Risk ownership makes AI risks accountable to governance bodies.
- **EU Trustworthy AI** — Technical robustness and safety: Risk-based approach. HLEG expects proportionate risk identification and mitigation.
- **AI Act** — Risk management: Risk management system. Direct mapping to AI Act risk management theme.
- **Mittelstadt et al.** — Inconclusive risk: Surface structural harms. Unowned risks remain inconclusive for affected communities.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

Operational Management (operational)

16. operational_monitoring

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Monitoring tracks performance, safety signals, and resource use.

Criterion description: Monitoring tracks performance, safety signals, and resource use.

Documentation hint (for field study context only; not rated in Round 1):

Monitoring dashboards, alert thresholds, and sample review reports.

Risk if missing (indicative):

Degradation, abuse, or cost overruns may persist without timely detection.

Traceability references:

- **ART** — Accountability: Operational performance and safety monitoring. Monitoring enables accountable response to degradation or misuse.
- **EU Trustworthy AI** — Technical robustness and safety: Reliability monitoring. Continuous monitoring supports robustness claims.
- **AI Act** — Post-market monitoring: Post-market monitoring. Operational metrics feed deployer monitoring duties.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

17. operational_incident_response

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Incident response covers harmful outputs, outages, and security events.

Criterion description: Incident response covers harmful outputs, outages, and security events.

Documentation hint (for field study context only; not rated in Round 1):

Incident response plan, tabletop exercise records, and post-incident review template.

Risk if missing (indicative):

Slow or inconsistent response can amplify harm to citizens and staff.

Traceability references:

- **ART** — Responsibility: Respond to harmful outputs and outages. Incident response assigns responsibility for remediation.
- **EU Trustworthy AI** — Technical robustness and safety: Serious incident handling. Preparedness aligns with safety and robustness expectations.
- **AI Act** — Post-market monitoring: Serious incidents. Reporting and response linked to post-market monitoring themes.
- **Mittelstadt et al.** — Unfair outcomes: Mitigate unfair harm. Slow response can allow unfair harms to citizens to accumulate.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

18. operational_human_oversight

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Human oversight is defined where outputs influence decisions or communications.

Criterion description: Human oversight is defined where outputs influence decisions or communications.

Documentation hint (for field study context only; not rated in Round 1):

Oversight procedures, sampling protocols, and appeal or correction pathways.

Risk if missing (indicative):

Automation bias and unreviewed errors may reach citizens or caseworkers.

Traceability references:

- **ART** — Responsibility: Human review of consequential outputs. Oversight assigns humans responsibility for consequential decisions.
- **Meaningful Human Control** — Human-in-the-loop / on-the-loop: Meaningful human intervention. Criterion operationalises meaningful human control for local LLM outputs.
- **EU Trustworthy AI** — Human agency and oversight: Human oversight. Direct HLEG requirement for human agency and oversight.
- **AI Act** — Human oversight: Human oversight measures. Maps to AI Act human oversight obligations for deployers.
- **Mittelstadt et al.** — Misguided agency: Counter automation bias. Absent oversight risks misguided agency where users over-trust automation.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):**Comment (optional):**

19. operational_documentation

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Documentation covers prompts, configurations, known limitations, and user guidance.

Criterion description: Documentation covers prompts, configurations, known limitations, and user guidance.

Documentation hint (for field study context only; not rated in Round 1):

System documentation set, prompt registry, and operator handbook.

Risk if missing (indicative):

Knowledge loss and inconsistent operation across shifts or sites.

Traceability references:

- **ART** — Transparency: Maintain prompts and limitations. Documentation makes system behaviour transparent to operators.
- **EU Trustworthy AI** — Transparency: Explainability and communication. Internal transparency supports trustworthy communication.
- **AI Act** — Transparency: Instructions for use. Operator documentation parallels transparency to deployers and staff.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):**Comment (optional):**

20. operational_lifecycle_management

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Lifecycle management covers pilot, production, scaling, and decommissioning.

Criterion description: Lifecycle management covers pilot, production, scaling, and decommissioning.

Documentation hint (for field study context only; not rated in Round 1):

Lifecycle stage gates, decommission checklist, and archival policy.

Risk if missing (indicative):

Experimental systems may remain in de facto production without controls.

Traceability references:

- **ART** — Accountability: Stage-gate pilot to decommission. Lifecycle gates maintain accountability as systems mature.
- **EU Trustworthy AI** — Technical robustness and safety: Lifecycle governance. Controlled transitions reduce safety gaps.
- **AI Act** — Risk management: Lifecycle risk management. Risk reassessment across lifecycle stages.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

Strategic Sovereignty (strategic_sovereignty)

21. strategic_sovereignty_vendor_independence

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Vendor independence reduces reliance on single proprietary stacks.

Criterion description: Vendor independence reduces reliance on single proprietary stacks.

Documentation hint (for field study context only; not rated in Round 1):

Alternative supplier analysis, open standards usage, and contingency planning.

Risk if missing (indicative):

Negotiating power and optionality decline as dependencies deepen.

Traceability references:

- **ART** — Responsibility: Reduce single-vendor dependence. Strategic responsibility for long-term optionality.
- **EU Trustworthy AI** — Diversity, non-discrimination and fairness: Diversity of supply. Vendor diversity mitigates structural dependency risks.
- **AI Act** — Risk management: Supply chain resilience. Third-party risk in deployer governance.
- **Mittelstadt et al.** — Unfair outcomes: Avoid coercive lock-in. Dependency can produce unfair bargaining outcomes for municipalities.

Your ratings

Field	Your response
-------	---------------

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

22. strategic_sovereignty_data_sovereignty

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Data sovereignty ensures primary data and models remain under institutional control.

Criterion description: Data sovereignty ensures primary data and models remain under institutional control.

Documentation hint (for field study context only; not rated in Round 1):

Data residency statement, on-prem storage design, and third-party access restrictions.

Risk if missing (indicative):

Sensitive public sector data may be exposed to external inference or training pipelines.

Traceability references:

- **ART** — Accountability: Institutional control of data and models. Sovereignty claims require demonstrable control over data assets.
- **EU Trustworthy AI** — Privacy and data governance: Privacy and data governance. Data sovereignty supports privacy-by-design for local LLMs.
- **AI Act** — Data governance: Data governance. Aligns with deployer data governance expectations.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

23. strategic_sovereignty_infrastructure_control

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Infrastructure control covers hosting, networking, and capacity planning.

Criterion description: Infrastructure control covers hosting, networking, and capacity planning.

Documentation hint (for field study context only; not rated in Round 1):

Infrastructure ownership model, capacity plan, and disaster recovery design.

Risk if missing (indicative):

Service continuity and security depend on opaque external platforms.

Traceability references:

- **ART** — Responsibility: Own hosting and capacity. Infrastructure control is a strategic responsibility for continuity.
- **EU Trustworthy AI** — Technical robustness and safety: Resilience. Resilient infrastructure underpins trustworthy operation.

- **AI Act** — Risk management: Operational resilience. Business continuity in risk management programmes.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

24. strategic_sovereignty_portability

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Portability supports migration of models, prompts, and evaluation assets.

Criterion description: Portability supports migration of models, prompts, and evaluation assets.

Documentation hint (for field study context only; not rated in Round 1):

Export formats, container images, and migration test results.

Risk if missing (indicative):

Switching costs may block adoption of improved local models or policies.

Traceability references:

- **ART** — Transparency: Migrate models and prompts. Portability makes technical commitments auditable and reversible.
- **EU Trustworthy AI** — Human agency and oversight: Reversibility and contestability. Portability supports contestability when systems fail citizens.
- **AI Act** — Risk management: Interoperability. Reduces dependency risk in technical governance.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

25. strategic_sovereignty_maintainability

Assessment question: For a bounded municipal AI/LLM programme dossier, is the following criterion relevant and clearly worded for programme-level governance readiness assessment? Criterion: Long-term maintainability plans staffing, upgrades, and technical debt management.

Criterion description: Long-term maintainability plans staffing, upgrades, and technical debt management.

Documentation hint (for field study context only; not rated in Round 1):

Roadmap, budget lines, and skills development plan for sustaining on-prem AI.

Risk if missing (indicative):

Systems may become unsupported while still handling citizen-facing workloads.

Traceability references:

- **ART** — Responsibility: Sustain staffing and upgrades. Long-term maintainability is a governance responsibility, not only IT.
- **EU Trustworthy AI** — Societal and environmental wellbeing: Long-term societal wellbeing. Sustainable operation aligns with HLEG societal wellbeing dimension.
- **AI Act** — Risk management: Lifecycle resource planning. Resource planning for ongoing compliance and monitoring.

Your ratings

Field	Your response
Relevance (1–5)	
Clarity (1–5)	
Essential? (Yes/No)	

Suggested revision (optional):

Comment (optional):

End of questionnaire — LocalGovBench v0.1 Delphi Round 1 (programme dossier / evidence-gated validation study)